Q: What hashing algorithm is used to detect unchanged documents in the ingestion pipeline?
A: SHA-256

Q: In the ingestion pipeline, which step is responsible for OCR and table preservation?
A: Document parsing & OCR (Step A2)

Q: What metadata is stored when a document is ingested?
A: Source URL, crawl timestamp, file hash, and language.

Q: How does the system avoid reprocessing files that haven't changed?
A: By hashing files (using SHA-256) and skipping content whose hash is unchanged.

Q: Why is semantic chunking preferred over fixed-size chunking in this RAG design?
A: Because it creates meaning-aligned retrieval units based on headings or structure rather than arbitrary token limits, improving retrieval relevance.

Q: Which pipeline step creates the single source of truth for documents, and what does it contain?
A: The Canonical Document Model (Step A3), containing a document ID, metadata, and ordered sections.

Q: What are the two key steps involved in the agentic RAG loop described?
A: User query rewriting and iterative self-evaluation by the LLM to determine if the answer is sufficient.

Q: Does the ingestion pipeline describe real-time document updates?
A: No, it explicitly describes an offline, resumable ingestion process.

Q: Which step removes headers and footers from documents?
A: Document parsing & OCR (Step A2).

Q: What is the main goal of the semantic chunking step?
A: To produce meaningful, retrieval-optimized units aligned to document structure and semantics.